

A Crash Course on SimpleCV

Katherine Scott

SightMachine

kat@sightmachine.com anthony@sightmachine.com

March 2, 2013

Overview

Quick Start!

What is SimpleCV?

What makes up SimpleCV?

SimpleCV

Getting Started

SimpleCV Shell

iPython Web Notebook

Image Basics

Really Basic Operations

Getting at the Pixels

Basic Manipulations

Let's talk about black, white, and color

Get Started!

There are a lot of dependencies for SimpleCV and it is a bit tough for beginners. We've brought disks that are ready to go!

- ▶ Windows / Linux
 - ▶ Boot from USB drive.
 - ▶ Alternatively install VirtualBox and the image.
 - ▶ <https://www.virtualbox.org/>
- ▶ Macs
 - ▶ Newer macs are persnickety about booting from a USB drive.
 - ▶ Install virtual box and the ISO and go to town.
- ▶ When you get home install from SuperPack or preferably source libs.
 - ▶ take awhile and is not a perfect science.
 - ▶ <https://github.com/ingenuitas/SimpleCV>
 - ▶ If you want to contribute this is a great place to start.



About the tutorial

- ▶ It will be a lot of live coding. I'll lead, you follow along.
- ▶ If you have a question feel free to interrupt.
- ▶ If you are having an issue raise a flag. Anthony will help you.

What makes up SimpleCV?

What Makes Up SimpleCV?



What makes up SimpleCV?

SimpleCV != OpenCV



- ▶ OpenCV is really busy, we help by wrapping python.
- ▶ We add lots of other fun stuff (OCR, Barcodes, etc.)
- ▶ We are not competing, we are complementing.
- ▶ Purposes are different. Python is great for prototyping. C++ great for embedded.

What makes up SimpleCV?

Core Dependencies

- ▶ OpenCV Python Bindings
- ▶ Numpy
- ▶ SciPy
- ▶ SciKits Learn and Orange
- ▶ PyGame (this is going away)
- ▶ Python Imaging Library (PIL)
- ▶ ipython
- ▶ PIL (Python Imaging Library)

What makes up SimpleCV?

Optional Dependencies

- ▶ Barcodes- Zebra Crossing ZXIng
- ▶ Optical Character Recognition (OCR) - Tesseract
- ▶ Beautiful Soup
- ▶ Kinect Support - freenect
- ▶ Unit Tests - nose
- ▶ Web Stuff - flask / CherryPy
- ▶ Arduino - pyfirmata
- ▶ Many Many Many more.

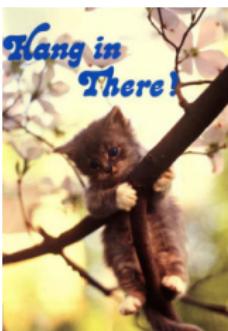
What makes up SimpleCV?

This is why we put everything in a superpack / virtual box / bootable drive

- ▶ Just get to the core library functions.
- ▶ We encourage you to install the full library when you get home.
- ▶ Help is available if you need it.

What makes up SimpleCV?

Getting Help after the tutorial.



- ▶ Primary Source: <http://help.simplecv.org/questions/>
- ▶ Documentation <http://www.simplecv.org/docs/>
- ▶ Tweet at us: @Simple_CV
- ▶ Another Good Resource:
<http://www.reddit.com/r/ComputerVision>

What makes up SimpleCV?

On the Printed Page

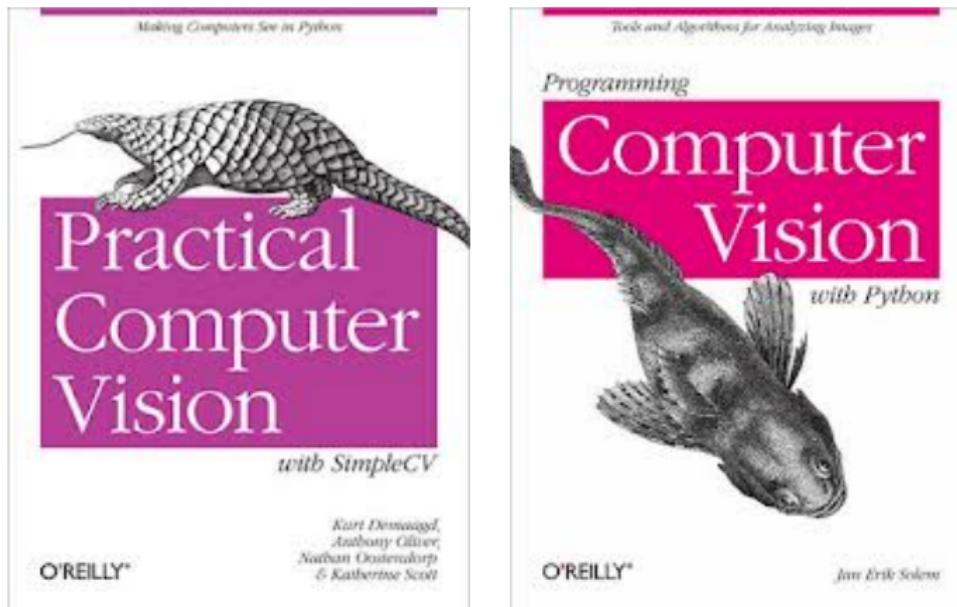


Figure: Two books about using Python for Computer Vision

What makes up SimpleCV?

So why are we doing this?

- ▶ We are really nice people who believe in Python and Open Source.
- ▶ We are trying to disrupt industrial quality control systems.



What makes up SimpleCV?

Early Prototypes

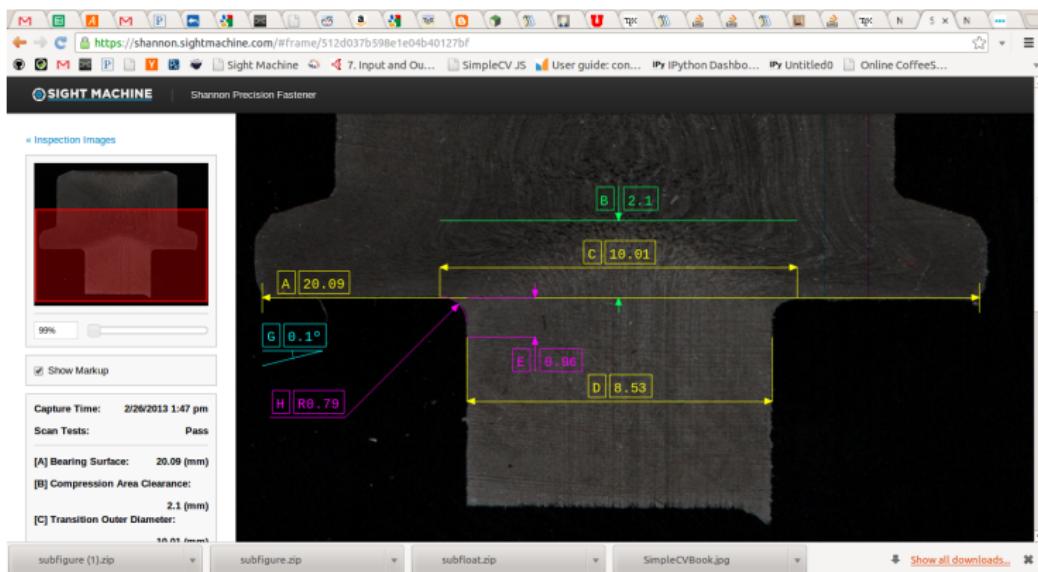


Figure: Early Customer - Industrial Fastener Morphology and Metallurgy

What makes up SimpleCV?

Early Prototypes

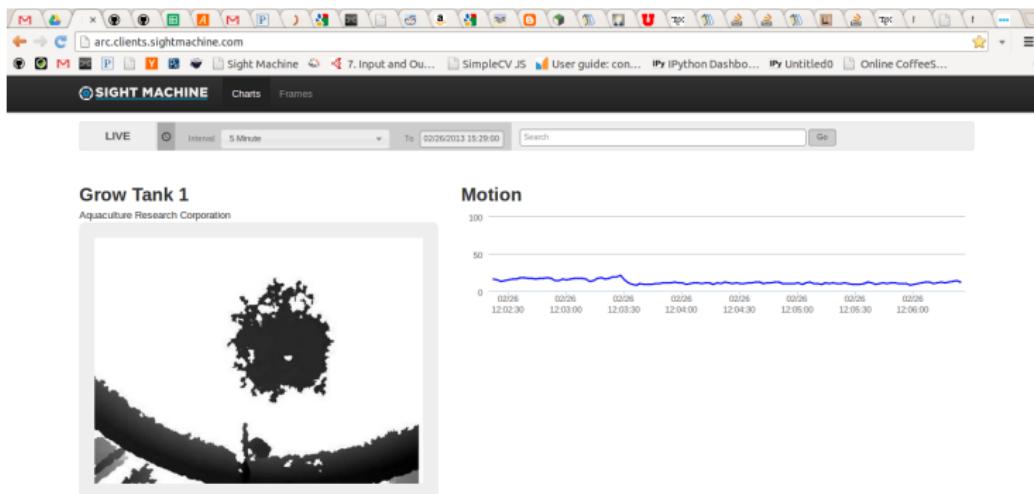
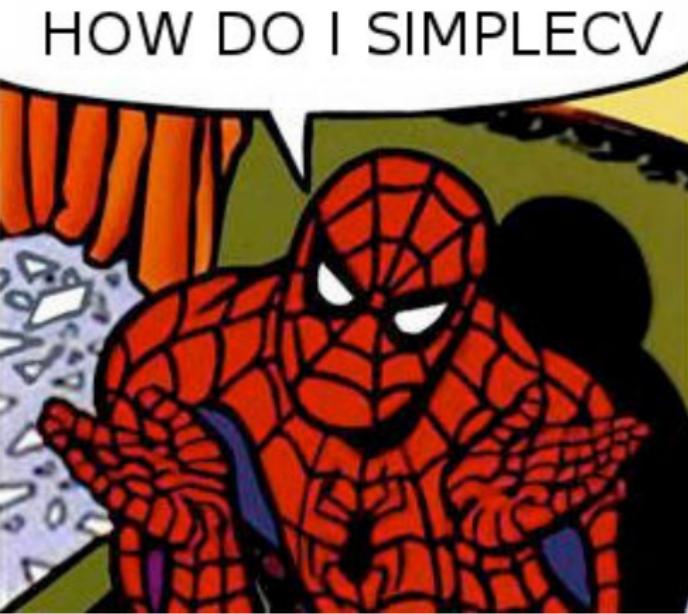


Figure: Early Customer - Aquaponics Research Facility

How do I SimpleCV?



Where do I write my code?

So how do I SimpleCV?

- ▶ In a python file, just like any other library.
- ▶ In a command line REPL like iPython.
- ▶ In the browser using iPython Notebooks (we'll use this today).

We really like iPython. It is kinda like using Matlab without the \$ 5000 per seat license cost.

Getting Started

How does fit into a work flow?

At SightMachine we roughly use these three tools for different parts of our workflow.

Tool	Uses
iPython REPL	Prototypes, Sanity Checks, Etc
iPython Web Notebook	Testing and Development
Python Files	Deployment Code

Table: SimpleCV Workflow

SimpleCV Hello World as a Script

Example (HelloWorld.py)

```
1 from SimpleCV import Image, Display, Color, Camera
2 cam = Camera(0) #Get the first camera
3 disp = Display((640,480)) # Create a 640x480 Display
4 while( disp.isNotDone() ):
5     img = cam.getImage() # get an image
6     # write text at 40,40 font_size 60pts, color is red
7     img.drawText("Hello World!",40,40,
8                 fontsize=60,color=Color.RED )
9     img.save(disp) # show it
10
```



Getting Started

How do I run Hello World?

- ▶ Run the py file with *python HelloWorld.py* in the command.
- ▶ Close it by pressing *esc* or *ctrl - c*

The image shows a Mac OS X desktop environment. On the left, a terminal window titled "Untitled window" displays a log of libusb debug messages. On the right, a video feed from a camera is displayed, showing a young girl with blonde hair giving a thumbs-up. The video frame has a red border. The desktop background is dark, and there are some icons at the bottom right.

```
libusb: 0.387092 debug [sysfs_scan_device] scan 2-1.1
libusb: 0.387131 debug [sysfs_scan_device] bus=2 dev=8
libusb: 0.387143 debug [enumerate_device] busnum 2 devaddr 8 session_id 520
libusb: 0.387156 debug [enumerate_device] allocating new device for 2/8 (session 520)
libusb: 0.387188 debug [sysfs_scan_device] scan 2-1.2
libusb: 0.387226 debug [sysfs_scan_device] bus=2 dev=8
libusb: 0.387239 debug [enumerate_device] busnum 2 devaddr 8 session_id 520
libusb: 0.387252 debug [enumerate_device]
libusb: 0.387256 debug [sysfs_scan_device]
libusb: 0.387325 debug [sysfs_scan_device]
libusb: 0.387358 debug [enumerate_device]
libusb: 0.387384 debug [sysfs_scan_device]
libusb: 0.387424 debug [sysfs_scan_device]
libusb: 0.387437 debug [enumerate_device]
libusb: 0.387449 debug [enumerate_device]
libusb: 0.387482 debug [discovered_device]
libusb: 0.387495 debug [sysfs_scan_device]
libusb: 0.387533 debug [sysfs_scan_device]
libusb: 0.387546 debug [enumerate_device]
libusb: 0.387559 debug [enumerate_device]
libusb: 0.387594 debug [libusb_get_device]
libusb: 0.387611 debug [libusb_get_device]
libusb: 0.387626 debug [libusb_get_device]
libusb: 0.387640 debug [libusb_get_device]
libusb: 0.387654 debug [libusb_get_device]
libusb: 0.387669 debug [libusb_get_device]
libusb: 0.387683 debug [libusb_get_device]
libusb: 0.387697 debug [libusb_get_device]
libusb: 0.387712 debug [libusb_get_device]
libusb: 0.387726 debug [libusb_get_device]
VIDIOC_QUERYMENU: Invalid argument
```

The SimpleCV Shell - Custom iPython REPL

Sometimes you just want to test an idea without writing a full script. For this reason we created the SimpleCV shell, which is a custom ipython instance. The SimpleCV shell will allow you to:

- ▶ Test your ideas in a REPL similar to Matlab.
- ▶ Access the SimpleCV documentation.
- ▶ Import modules that you are working with to test.
- ▶ Run through an interactive tutorial.



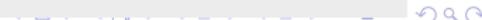
SimpleCV Shell

Starting the SimpleCV Shell

In OSX and Linux just type *simplecv* at the command line. On Windows you just click on the SimpleCV icon.

Example (Shell Basics)

```
+-----+
SimpleCV 1.3.0 [interactive shell] - http://simplecv.org
+-----+
Commands:
"exit()" or press "Ctrl+ D" to exit the shell
"clear" to clear the shell screen
"tutorial" to begin the SimpleCV interactive tutorial
"example" gives a list of examples you can run
"forums" will launch a web browser for the help forums
"walkthrough" will launch a web browser with a walkthrough
Usage:
dot complete works to show library
for example: Image().save("/tmp/test.jpg") will dot complete
just by touching TAB after typing Image().
Documentation:
help(Image), ? Image, Image ?, or Image() ? all do the same
"docs" will launch webbrowser showing documentation
```



SimpleCV Shell Like a Boss



- ▶ Putting a ? in front of a class or method will give you documentation. The "/" key will let you search.
- ▶ iPython has tab completion for methods.
- ▶ Up arrow will give you previous commands.
- ▶ %paste will let you paste formatted code.
- ▶ Other cool stuff can be found by googling iPython magic

SimpleCV Shell

Let's repeat Hello World in SimpleCV Shell

Example (In the SimpleCV shell)

```
SimpleCV:1> cam = Camera()
SimpleCV:2> disp = Display((640,480))
SimpleCV:3> while disp.isNotDone():
...:     img = cam.getImage().edges()
...:     img.drawText("Hello World!",40,40,fontsize=60)
...:     img.save(disp)
...:
SimpleCV:4> exit
```

- ▶ Just push return after each line.
- ▶ iPython will do tabbing in the while loop.
- ▶ *esc* to quit or *ctrl - c*.
- ▶ type “exit” to quit.

SimpleCV Shell

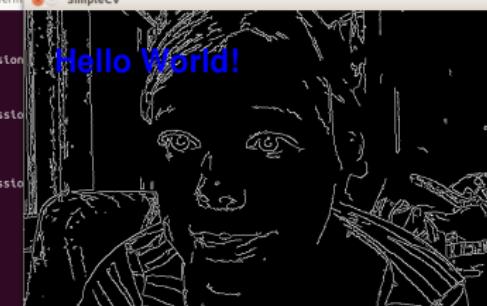
Yes, it really is that simple.

```
pygame window Terminal Term SimpleCV
[libusb: 3.5087961 debug [sysfs_scan_device] scan 2-1.1.1
[libusb: 3.5088722 debug [sysfs_scan_device] bus=2 dev=9
[libusb: 3.5088989 debug [enumerate_device] busnum 2 devaddr 9 session_id 521
[libusb: 3.5081044 debug [enumerate_device] allocating new device for 2/9 (session
[libusb: 3.5081933 debug [sysfs_scan_device] scan 2-1.1.2
[libusb: 3.5083022 debug [sysfs_scan_device] bus=2 dev=10
[libusb: 3.5083194 debug [enumerate_device] busnum 2 devaddr 10 session_id 522
[libusb: 3.5083344 debug [enumerate_device] allocating new device for 2/10 (session
[libusb: 3.5084211 debug [discovered_devs_append] need to increase capacity
[libusb: 3.5084411 debug [sysfs_scan_device] scan 2-1.1.3
[libusb: 3.5085595 debug [sysfs_scan_device] bus=2 dev=11
[libusb: 3.5085771 debug [enumerate_device] busnum 2 devaddr 11 session_id 523
[libusb: 3.5085922 debug [enumerate_device] allocating new device for 2/11 (session
[libusb: 3.5087444 debug [libusb_get_device_descriptor]
[libusb: 3.5087833 debug [libusb_get_device_descriptor]
[libusb: 3.5088133 debug [libusb_get_device_descriptor]
[libusb: 3.5088444 debug [libusb_get_device_descriptor]
[libusb: 3.5088744 debug [libusb_get_device_descriptor]
[libusb: 3.5089655 debug [libusb_get_device_descriptor]
[libusb: 3.5089371 debug [libusb_get_device_descriptor]
[libusb: 3.5089671 debug [libusb_get_device_descriptor]
[libusb: 3.5089991 debug [libusb_get_device_descriptor]
[libusb: 3.5090381 debug [libusb_get_device_descriptor]

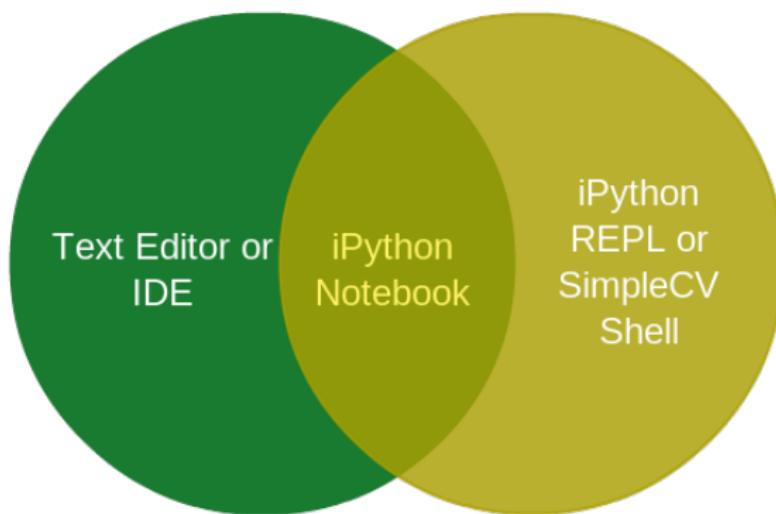
VIDEOC_QUERYMENU: Invalid argument

SimpleCV:2> disp = Display((640,480))

SimpleCV:3> while disp.isNotDone():
...     img = cam.getImage().edges()
...     img.drawText('Hello World!', 40,40,fontsize=60)
...     img.save(disp)
... 
```



Why use iPython Web Notebooks

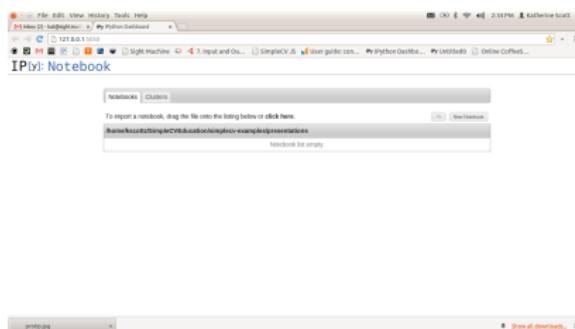


[online diagramming & design]  .com

- ▶ Web notebooks give you the best features of an IDE and a REPL



How do I use the notebook?



- ▶ From the shell just type *simplecv notebook -- pylab inline*.
- ▶ The *-- pylab inline* is optional but it pulls in matplotlib which is handy.
- ▶ You will get to a dashboard to create a new notebook.
- ▶ By default notebooks are in the path where you start ipython.





iPython Web Notebook

How do I use the notebook?

The screenshot shows an iPython Notebook window titled "IP[y]: Notebook" with the URL "127.0.0.1:5050/256fbf32-163e-433b-8621-293fe3f33e48". The notebook has two cells:

- In [1]:** `Image?`
- In [2]:** (empty)

The output for In [1] is:

```
Type: classobj
String Form:SimpleCV.ImageClass.Image
File: /home/kscott/SimpleCV/SimpleCV/ImageClass.py
Docstring:
+SUMMARY+
The Image class is the heart of SimpleCV and allows you to convert to and from a number of source types with ease. It also has intelligent buffer management, so that modified copies of the Image required for algorithms such as edge detection, etc can be cached and reused when appropriate.

Image are converted into 8-bit, 3-channel images in RGB colorspace. It will automatically handle conversion from other representations into this standard format. If dimensions are passed, an empty image is created.

+EXAMPLE+
>>> i = Image('/path/to/image.png')
>>> i = Camera().getImage()
```

At the bottom, there is a "proto.ipynb" button and a "Show all downloads..." link.

- ▶ Everything we mentioned about the SimpleCV shell still holds.
- ▶ Magic commands, inline documentation, etc. still work.
- ▶ *enter* starts a new line.
- ▶ *ctrl – enter* executes a line.



Caveats about iPython Web Notebooks



- ▶ iPython Web Notebooks are still version 0.1.4
- ▶ **There is no auto-save. Get in the `ctrl - s` save habit.**
- ▶ If you edit a module you import you must restart the core.
- ▶ Minimal editing support. No find/replace.
- ▶ The core can sometimes crash on large images.
- ▶ The notebooks hold on to data by default. This can fill up your version control system fast. Try the download as python command from the gui.

Image Loading Basics II

- ▶ You can get the image file name using `img.filename`
- ▶ Images can also come from appropriately shaped numpy arrays.
- ▶ PIL and OpenCV images can also be passed into the image.
- ▶ Can also take a URL to an image.
- ▶ The `img.getEXIFData()` command can show jpg EXIF data.

Saving an Image

- ▶ The `img.save()` command is used to save images.
- ▶ You can save as just about any format, PNG, JPG, WebP, etc.
- ▶ Calling `save` with no parameters saves it to temp directory.
- ▶ Using the `params` flag you can set compression, e.g. set compression quality.

Using a USB Camera

- ▶ Most usb cameras use the **Camera** class.
- ▶ The camera class takes a camera index, *usually* this is the order cameras are plugged into the computer.
- ▶ The camera also has properties that you can get and set, or use a propmap dictionary to set.
- ▶ *Support for camera properties is vendor specific and spotty at best.*
- ▶ Cameras also have a *threaded* parameter. Set this to false to run multiple cameras.

oooooooooooo

ooooooo
ooooo
oooo
oooo

oooooo

oooooooooooo

ooooooo
ooooo
oooo
oooo

oooooo

- ▶ Camera.getImage() will return the current image.
- ▶ Camera.getAllProperties() will return the cameras properties.
- ▶ Camera.getProperty() and Camera.setProperty() may let you set properties. *This is highly vendor dependent and usually poorly documented.*

oooooooooooo

ooooooo
ooooo
oooo
oooo

oooooo

- ▶ Camera.getImage() will return the current image.
- ▶ Camera.getAllProperties() will return the cameras properties.
- ▶ **Camera.getProperty() and Camera.setProperty()** may let you set properties. *This is highly vendor dependent and usually poorly documented.*

Briefly: Other Cameras

- ▶ Kinect - Depth Camera
 - ▶ Uses freenect drivers, not the OpenNI drivers.
 - ▶ `Kinect.getImage` and `Kinect.getDepth`
 - ▶ Note that these aren't well calibrated together.
- ▶ JpegStreamReader - IP Cameras
 - ▶ Give it a url to camera's web feed, and scrape images.
 - ▶ Getting the URL straight can be tricky.
- ▶ Virtual Camera
 - ▶ A virtual camera that pulls from a directory full of images, or video.
 - ▶ Interface to video files for processing.

Briefly: Other Cameras

- ▶ Document Scanners
 - ▶ SANE compatible devices.
 - ▶ Allow you to set resolution and ROI.
- ▶ Digital Camera
 - ▶ Uses Piggy Photo Library
 - ▶ Works with most DSLRs and point and shoots.
- ▶ AVT Camera
 - ▶ Professional digital imaging cameras with interchangeable optics.
 - ▶ Fine grain control of camera parameters.

Image Sets

ImageSets are lists of images. They are great for aggregating datasets.

- ▶ By default load all image files in a directory.
- ▶ Can iterate over the list using list comps or for loops.
- ▶ Using the BeautifulSoup library can download sets from google.
- ▶ Can save image sets to directories, or animated gifs!
- ▶ The show command works just like on image class.
- ▶ Can apply averages to images.
- ▶ More coming soon.

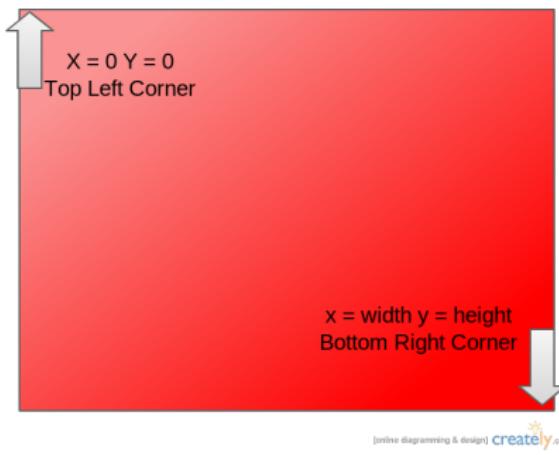
ImageSet Example

Example (Image Sets)

```
mySet = ImageSet()  
mySet.download('cats') # download cats  
mySet.show() # show them to us  
avgCat = mySet.average() # get the avg cat  
avgCat.show() # show the avg cat  
mySet[3].show() # show the third kitty  
resized = mySet.standardize(128,64)  
resized.save('cat.gif')  
mySet.save('cat.gif') # save the cats as a gif  
for cat in mySet: # iterate over the set of cats  
    cat.binarize().show()
```

Getting at the Pixels

Getting at a pixel



- ▶ SimpleCV treats images as two dimensional arrays of color value tuples.
- ▶ Each tuple holds three values (Red,Green,Blue).
- ▶ Pixels start in the top left corner at zero.

Getting at the Pixels

Pixel Manipulation Example

Example (Getting at those pixels)

```
img = Image('helloworld.jpg')
c = img[0,0] # get a pixel
print c
print img.getPixel(0,0) # get another way
test = img[200:300,200:300]
test.show() # the result is an image
test = img[50,:,:]
test.show() # again using slice
test = img[0:5,0:5]
print test.getNumpy() # get the raw values
img[0:105,0:105] = Color.RED
img.show() # Another way
test = img.getNumpy() #RIGHT!
test[0:105,0:105] = Color.RED
img2 = Image(test)
img2.show()
```

Getting at the Pixels

Getting at a pixel

- ▶ Images are read only. To write a pixel directly you need to create a new image. Usually this happens in numpy
- ▶ Images support the list slice notation but return images.
- ▶ To get at the raw pixel values use **getNumpy()** and **getGrayNumpy()**
- ▶ Numpy values can be accessed using slices the first parameter is x, the second is y, and the third is the channel in RGB order. For example `npimg[x][y][0]`.
- ▶ The **Image.width** and **Image.height** member variables can help you find your way.

Another Pixel Manipulation Example

Example (Fancy Manipulations)

```
img = Image('helloworld.jpg')
gray = img.getGrayNumpy() # get the gray scale np image
colored = img.getNumpy() # and the colored one
print (img.width,img.height) #tell us the image size
colored[0:20,:] = Color.BLUE # set the left side to blue
colored[:,0:20] = Color.GREEN # set the top row to green
colored[40:80,40:80] \quad
    \includegraphics[width=0.4\linewidth]{JanEricBook.jpg}
= Color.YELLOW # make a yellow square
x,y = np.where(gray>230) # find bright pixels > 230
for xf,yf in zip(x,y): # for each of those
    colored[xf][yf] = Color.RED #make them red
img2 = Image(colored) # create an image
img2.show() # and show it
# now set the whole blue channel to 255
colored[:, :, 2] = 255
# and show us that
img3 = Image(colored)
img3.show()
```

Getting at the Pixels

Getting at a pixel

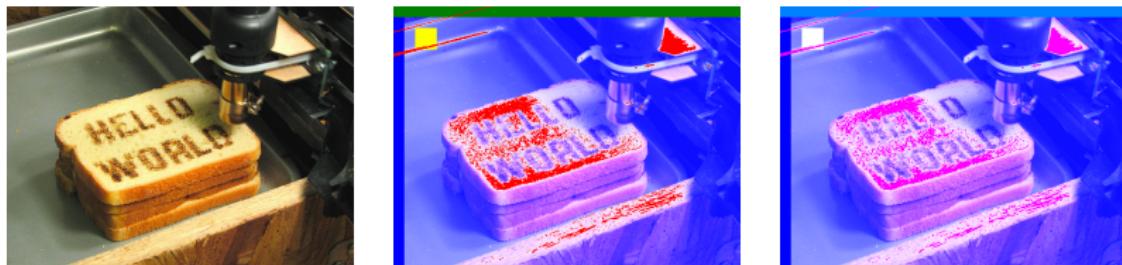


Figure: Not bad for less than 20 lines of code.

Cropping, Scaling, Rotating, etc

It is helpful to think of some image processing like using an image editor, like GIMP, paint, or that other one Adobe makes. Here are a few basic image operations you can do in SimpleCV.

- ▶ **Image.crop**. Does what it says on the tin. We refer to crop areas by their top left corner x and y plus the width height.
- ▶ **Image.scale** scales the image proportionally while **Image.resize** resizes the image to the desired size.
 - ▶ **Image.resize** is smart, if you tell it a width or height it will infer the other parameter from the aspect ratio.
 - ▶ **Image.scale** uses a proportionality. So passing it a value of two will double the image size. Make sure to check out the interpolation method.

Rotating

- ▶ **Image.rotate** takes and angle in degrees.
- ▶ Rotate has a parameter called fixed. If fixed is set to false, SimpleCV will create a new image that matches the rotated image size. Otherwise we stick the rotated image in data in a similarly sized canvas.
- ▶ It is possible to pick the x,y position of the rotation. The default is the center of the image.
- ▶ You can also scale an image while rotating.
- ▶ In a pinch you can use **Image.flipHorizontal** and **Image.flipVertical**.
- ▶ Be aware **FLIPPING != ROTATION**

Blit

The **Image.blit** function gets its name from the old computer graphics term “bit block image transfer”. Really it means just copy and paste another image onto another image and return the result.

- ▶ Blit takes in another image and a position where you want to put it.
- ▶ That position can be negative with respect to the destination image. For example (-10,-10) would put the source image over the top left of the destination image.
- ▶ You can toss blit a binary mask, an alpha value (that is transparency) or a an alpha mask (a grayscale image that has an alpha mask per pixel).

Let's play!

Example (Basic Manipulations)

```
img = Image('tricky.jpg')
face = img.crop(150,190,309-150,333-190)
img.show() # show the source
face.show() # show the cropped image
face.rotate(45).show() # basic rotation
face.rotate(45,fixed=False).show()
face.scale(0.5).show()
face.scale(width=int(face.width/2.0)) # basic scaling
face.flipHorizontal().show() # flipping
test1 = img.blit(face.flipHorizontal(),pos=(150,190))
test1.show() # now let's have fun with blitting
test2 = img.blit(face.flipHorizontal(),pos=(150,190),alpha=0.5)
test2.show()
mask = Image((face.width,face.height))
# Here we are just drawing a white circle on a black background
mask.drawCircle((face.width/2,face.height/2),70,color=Color.WHITE, thickness=-1)
mask = mask.applyLayers()
mask.show()
test3 = img.blit(face.flipHorizontal(),pos=(150,190),mask=mask)
test3.show()
face.binarize().blur().show()
test4 = img.blit(face.flipHorizontal(),pos=(150,190),alphaMask=face.binarize().blur())
test4.show()
```



The Blit Progression of President Dick Nixon



Figure: Various blitting

Can you make our subject a smaller face or a slightly slanted face?

oooooooooooo

ooooooo
ooooo
oooo
oooo

oooooo

A few more basics

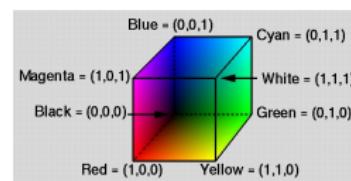
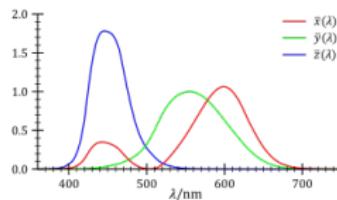
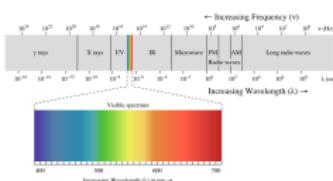
- ▶ **Image.blur** - A basic Gaussian blur.
- ▶ **Image.smooth** - A variety of smoothing filters. Some good for removing camera noise.
- ▶ **Image.toGray** - Convert a color image to a gray scale image.
- ▶ **Image.threshold** - Take an image and set all off the pixels that have a grayscale value above the threshold to white, and everything else to black.
- ▶ **Image.invert** - Swap the lightest and darkest values, like a photo negative.

Let's Talk about color, color spaces, and much more.



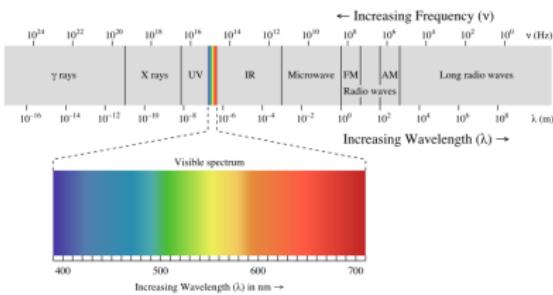
Computer vision is all about moving from a vast amount of information to small result as quickly as possible. Working in grayscale, or black and white images can speed things up dramatically.

Color is surprisingly hard to represent



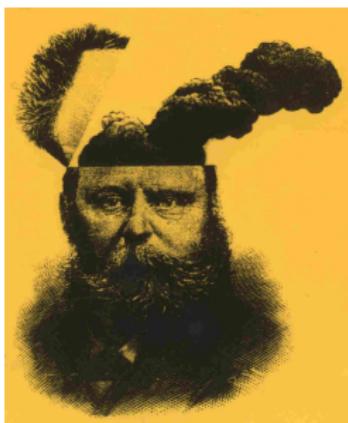
We start with visible light, which bounces around, does funky stuff, and then enters our eye or camera. Our eyes and cameras have a response function that does an imperfect job of sampling the parts of the spectrum. We then take those samples and try to map them onto finite color space like RGB.

An illustration of why color is hard.



Point to where this magenta flower lives on the visible light spectrum.

MIND BLOWN!



Magenta doesn't exist in nature. It is trick our brains and cameras play on us. It exists because we sample the spectrum and try to recombine the samples. We wrap the visible spectrum around.

oooooooooooo

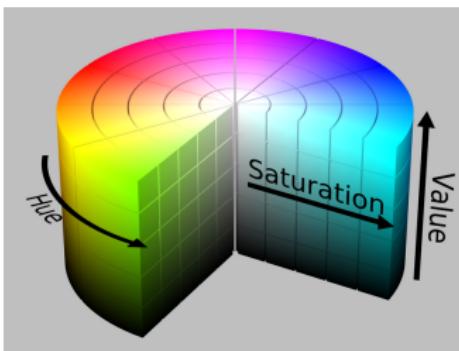
ooooooo

oooooo

oooo

oooo

To manage this problem we use color spaces.



- ▶ To deal with this problem we use color spaces.
- ▶ Most images are in the RGB, BGR, or gray scale color spaces.
- ▶ Sometimes it is helpful to use the hue, saturation, and value (HSV) color space. In HSV you have a
 - ▶ Hue, or pure color (from 0 to 180)
 - ▶ Saturation that tells us how far from white our color is
 - ▶ Value that tells us how dark the color is.

How do I work with Color in SimpleCV

- ▶ The **Image.toHSV()** function will convert your image, while the **Image.toBGR()** function will return it back to the original.
- ▶ We also have **Image.toGray()** and **Image.toRGB()** and whole bunch more.
- ▶ The **Image.huePeaks** function can help you guess what hues are in your image.
- ▶ The **Image.hueDistance** function can then help you look for stuff that is the hue you want.
- ▶ The **Image.hueHistogram** function will let you visualize the results.

Finding colors

Example (Using Hue)

```
import pylab as plt # import pylab for plotting
img = Image("flower.jpg")
img.show()
peaks = img.huePeaks() # find the hues
print peaks
hist = img.hueHistogram() # get the histogram too
plt.plot(hist) # plot the histogram
plt.draw() # show it to us
# show how far away from the hue we are black is closer
# white is farther away from our hue.
hdist = img.hueDistance(peaks[3][0])
# create a black and white version of our flower
binary = img.hueDistance(peaks[3][0]).invert().threshold(220)
hdist.show()
binary.show()
hdist.save('hflower.png')
binary.save('bflower.png')
```

oooooooooooo

ooooooo

oooooo

ooooo

oooo

oooo

Color Finding Results

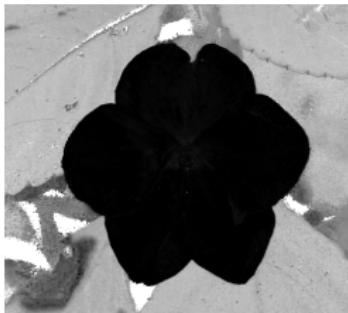


Figure: Finding an object from colors

Paragraphs of Text

Sed iaculis dapibus gravida. Morbi sed tortor erat, nec interdum arcu. Sed id lorem lectus. Quisque viverra augue id sem ornare non aliquam nibh tristique. Aenean in ligula nisl. Nulla sed tellus ipsum. Donec vestibulum ligula non lorem vulputate fermentum accumsan neque mollis.

Sed diam enim, sagittis nec condimentum sit amet, ullamcorper sit amet libero. Aliquam vel dui orci, a porta odio. Nullam id suscipit ipsum. Aenean lobortis commodo seDerkt commodo leo gravida vitae. Pellentesque vehicula ante iaculis arcu pretium rutrum eget sit amet purus. Integer ornare nulla quis neque ultrices lobortis. Vestibulum ultrices tincidunt libero, quis commodo erat ullamcorper id.

Bullet Points

- ▶ Lorem ipsum dolor sit amet, consectetur adipiscing elit
- ▶ Aliquam blandit faucibus nisi, sit amet dapibus enim tempus eu
- ▶ Nulla commodo, erat quis gravida posuere, elit lacus lobortis est, quis porttitor odio mauris at libero
- ▶ Nam cursus est eget velit posuere pellentesque
- ▶ Vestibulum faucibus velit a augue condimentum quis convallis nulla gravida

Multiple Columns

Heading

1. Statement
2. Explanation
3. Example

*Lorem ipsum dolor sit amet,
consectetur adipiscing elit.
Integer lectus nisl, ultricies in
feugiat rutrum, porttitor sit amet
augue. Aliquam ut tortor mauris.
Sed volutpat ante purus, quis
accumsan dolor.*

Table

Treatments	Response 1	Response 2
Treatment 1	0.0003262	0.562
Treatment 2	0.0015681	0.910
Treatment 3	0.0009271	0.296

Table: Table caption

Theorem

Theorem (Mass–energy equivalence)

$$E = mc^2$$

Verbatim

Example (Theorem Slide Code)

```
def doStuff(a,b,c=[1,2,3]):  
    a = 5  
    b = a  
    c.reverse()  
  
derp = [1,2,3,4]  
for i in derp:  
    doStuff()  
    pass  
print deep
```

Verbatim

Example (Theorem Slide Code)

```
$E = mc^2$
```

Figure

Uncomment the code on this slide to include your own image from the same directory as the template .TeX file.

Citation

An example of the \cite command to cite within the presentation:

This statement requires citation [Smith, 2012].

References



John Smith (2012)

Title of the publication

Journal Name 12(3), 45 – 678.

The End